

## ASSIGNMENT 5.1

### Datasets used:

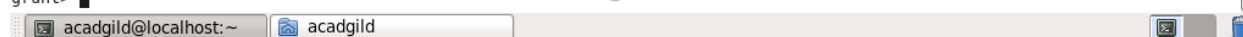
#### employee\_details.txt

```
employee_details = LOAD 'employee_details.txt' USING PigStorage(',') AS (emp_id:int,  
emp_name:chararray, emp_salary:int, emp_rating:int);
```

### Verification:

```
DUMP 'employee_details';
```

```
2017-12-10 19:36:12,566 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 19:36:12,566 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to pro
cess : 1
(101,Amitabh,20000,1)
(102,Shahrukh,10000,2)
(103,Akshay,11000,3)
(104,Anubhav,5000,4)
(105,Pawan,2500,5)
(106,Aamir,25000,1)
(107,Salman,17500,2)
(108,Ranbir,14000,3)
(109,Katrina,1000,4)
(110,Priyanka,2000,5)
(111,Tushar,500,1)
(112,Ajay,5000,2)
(113,Jubeen,1000,1)
(114,Madhuri,2000,2)
grunt> DESCRIBE employee_details;
employee_details: {emp_id: int,emp_name: chararray,emp_salary: int,emp_rating: int}
grunt> █
```



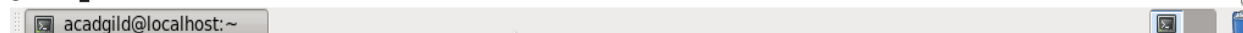
#### employee\_expenses.txt

```
employee_expenses = LOAD 'employee_expenses.txt' AS (emp_id:int, emp_expense:int);
```

### Verification:

```
DUMP 'employee_expenses';
```

```
2017-12-10 23:33:35,160 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 23:33:35,160 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to pro
cess : 1
(101,200)
(102,100)
(110,400)
(114,200)
(119,200)
(105,100)
(101,100)
(104,300)
(102,400)
grunt> DESCRIBE employee_expenses;
employee_expenses: {emp_id: int,emp_expense: int}
grunt> █
```



**Task 1:** Top 5 employees (employee id and employee name) with highest rating. (In case two employees have same rating, employee with name coming first in dictionary should get preference)

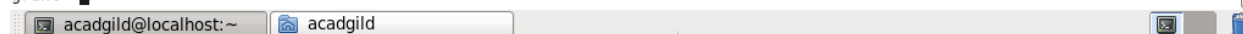
### Solution:

Step 1: Sort employees by their ratings in descending order.

```
sorted_employees = ORDER employee_details BY emp_rating DESC;
```

Output: DUMP sorted\_employees;

```
2017-12-10 19:48:48,726 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 19:48:48,726 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(105,Pawan,2500,5)
(110,Priyanka,2000,5)
(104,Anubhav,5000,4)
(109,Katrina,1000,4)
(108,Ranbir,14000,3)
(103,Akshay,11000,3)
(114,Madhuri,2000,2)
(112,Ajay,5000,2)
(107,Salman,17500,2)
(102,Shahrukh,10000,2)
(111,Tushar,500,1)
(113,Jubeen,1000,1)
(101,Amitabh,20000,1)
(106,Aamir,25000,1)
grunt> DESCRIBE sorted_employees;
sorted_employees: {emp_id: int,emp_name: chararray,emp_salary: int,emp_rating: int}
grunt>
```

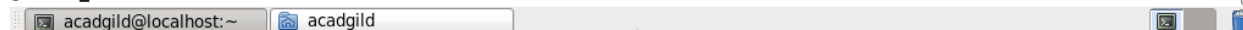


Step 2: Get top 5 employees.

```
top_5_employees = LIMIT sorted_employees 5;
```

Output: DUMP top\_5\_employees;

```
2017-12-10 19:49:46,525 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 19:49:46,526 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(110,Priyanka,2000,5)
(105,Pawan,2500,5)
(109,Katrina,1000,4)
(104,Anubhav,5000,4)
(108,Ranbir,14000,3)
grunt> DESCRIBE top_5_employees;
top_5_employees: {emp_id: int,emp_name: chararray,emp_salary: int,emp_rating: int}
grunt>
```



Step 3: Get employee id and employee name from above relation.

```
top_5_employees_final = FOREACH top_5_employees GENERATE emp_id, emp_name;
```

Output: DUMP top\_5\_employees\_final;

```
2017-12-10 19:50:46,719 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 19:50:46,719 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(110,Priyanka)
(105,Pawan)
(109,Katrina)
(104,Anubhav)
(108,Ranbir)
grunt> DESCRIBE top_5_employees_final;
top_5_employees_final: {emp_id: int,emp_name: chararray}
grunt>
```

**Task 2:** Top 3 employees (employee id and employee name) with highest salary, whose employee id is an odd number. (In case two employees have same salary, employee with name coming first in dictionary should get preference).

**Solution:**

Step 1: Filter out and capture employees with odd employee id.

employees\_with\_odd\_emp\_id = FILTER employee\_details BY emp\_id%2 == 1;

Output: DUMP employees\_with\_odd\_emp\_id;

```
2017-12-11 18:11:38,826 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-11 18:11:38,826 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh,20000,1)
(103,Akshay,11000,3)
(105,Pawan,2500,5)
(107,Salman,17500,2)
(109,Katrina,1000,4)
(111,Tushar,500,1)
(113,Jubeen,1000,1)
grunt> DESCRIBE employees_with_odd_emp_id;
employees_with_odd_emp_id: {emp_id: int,emp_name: chararray,emp_salary: int,emp_rating: int}
grunt>
```

Step 2: Sort employees by their salary in descending order.

sort\_by\_emp\_salary = ORDER employees\_with\_odd\_emp\_id BY emp\_salary DESC;

Output: DUMP sort\_by\_emp\_salary;

```
2017-12-11 18:14:03,615 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-11 18:14:03,615 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh,20000,1)
(107,Salman,17500,2)
(103,Akshay,11000,3)
(105,Pawan,2500,5)
(113,Jubeen,1000,1)
(109,Katrina,1000,4)
(111,Tushar,500,1)
grunt> DESCRIBE sort_by_emp_salary;
sort_by_emp_salary: {emp_id: int,emp_name: chararray,emp_salary: int,emp_rating: int}
grunt>
```

Step 3: Get top 3 employees.

```
top_3_emp_by_salary = LIMIT sort_by_emp_salary 3;
```

Output: DUMP top\_3\_emp\_by\_salary;

```
2017-12-11 18:15:11,042 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-11 18:15:11,042 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh,20000,1)
(107,Salman,17500,2)
(103,Akshay,11000,3)
grunt> DESCRIBE top_3_emp_by_salary;
top_3_emp_by_salary: {emp_id: int,emp_name: chararray,emp_salary: int,emp_rating: int}
grunt>
```

Step 4: Get employee id and employee name from above relation.

```
top_3_emp_by_salary_final = FOREACH top_3_employees GENERATE emp_id, emp_name;
```

Output: DUMP top\_3\_emp\_by\_salary\_final;

```
2017-12-11 18:09:20,141 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-11 18:09:20,141 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh)
(107,Salman)
(103,Akshay)
grunt> DESCRIBE top_3_emp_by_salary_final;
top_3_emp_by_salary_final: {emp_id: int,emp_name: chararray}
grunt>
```

**Task 3:** Employee (employee id and employee name) with maximum expense (In case two employees have same expense, employee with name coming first should get preference)

**Solution:**

Step 1: Join relations employee\_details and employee\_expenses since we require data from both.

```
joined_emp_info = JOIN employee_details BY emp_id, employee_expenses BY emp_id;
```

Output: DUMP joined\_emp\_info;

```
2017-12-10 23:24:42,092 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 23:24:42,092 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh,20000,1,101,100)
(101,Amitabh,20000,1,101,200)
(102,Shahrukh,10000,2,102,400)
(102,Shahrukh,10000,2,102,100)
(104,Anubhav,5000,4,104,300)
(105,Pawan,2500,5,105,100)
(110,Priyanka,2000,5,110,400)
(114,Madhuri,2000,2,114,200)
grunt> DESCRIBE joined_emp_info;
joined_emp_info: {employee_details::emp_id: int,employee_details::emp_name: chararray,employee_details::emp_salary: int,employee_details::emp_rating: int,employee_expenses::emp_id: int,employee_expenses::emp_expense: int}
grunt>
```

Step 2: Group data from joined relation by all columns i.e. each tuple/record will be put in separate group.

grouped\_emp\_info = GROUP joined\_emp\_info ALL;

Output: DUMP grouped\_emp\_info;

```
2017-12-10 23:25:49,102 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 23:25:49,102 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(all, {(114, Madhuri, 2000, 2, 114, 200), (110, Priyanka, 2000, 5, 110, 400), (105, Pawan, 2500, 5, 105, 100), (104, Anubhav, 5000, 4, 104, 300), (102, Shahrukh, 10000, 2, 102, 100), (102, Shahrukh, 10000, 2, 102, 400), (101, Amitabh, 20000, 1, 101, 200), (101, Amitabh, 20000, 1, 101, 100)})
grunt> DESCRIBE grouped_emp_info;
grouped_emp_info: {group: chararray, joined_emp_info: {(employee_details::emp_id: int, employee_details::emp_name: chararray, employee_details::emp_salary: int, employee_details::emp_rating: int, employee_expenses::emp_id: int, employee_expenses::emp_expense: int)}}
```

Step 3: Find out the maximum expense value.

max\_expense = FOREACH grouped\_emp\_info GENERATE MAX (joined\_emp\_info.  
employee\_expenses::emp\_expense) AS expense;

Output: DUMP max\_expense;

```
2017-12-10 23:26:57,300 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 23:26:57,300 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(400)
grunt> DESCRIBE max_expense;
max_expense: {expense: int}
grunt>
```

Step 4: Find out the employees who has got expense same as the maximum value.

emp\_with\_max\_expense = FILTER joined\_emp\_info BY employee\_expenses::emp\_expense ==  
max\_expense.expense;

Output: DUMP emp\_with\_max\_expense;

```
2017-12-10 23:28:37,360 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 23:28:37,361 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(102, Shahrukh, 10000, 2, 102, 400)
(110, Priyanka, 2000, 5, 110, 400)
grunt> DESCRIBE emp_with_max_expense;
emp_with_max_expense: {employee_details::emp_id: int, employee_details::emp_name: chararray, employee_details::emp_salary: int, employee_details::emp_rating: int, employee_expenses::emp_id: int, employee_expenses::emp_expense: int}
grunt>
```

Step 5: If there are two or more employees with maximum expenses, sort them out in ascending order.

sorted\_by\_emp\_name = ORDER emp\_with\_max\_expense BY emp\_name;

Output: DUMP sorted\_by\_emp\_name;

```

2017-12-10 23:29:17,826 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 23:29:17,826 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(110,Priyanka,2000,5,110,400)
(102,Shahrukh,10000,2,102,400)
grunt> DESCRIBE sorted_by_emp_name;
sorted_by_emp_name: {employee_details::emp_id: int,employee_details::emp_name: chararray,employee_details::emp_salary: int,employee_details::emp_rating: int,employee_expenses::emp_id: int,employee_expenses::emp_expense: int}
grunt>

```

acadgild@localhost: ~

Step 6: Take first employee record from the above result.

```
get_top_emp = LIMIT sorted_by_emp_name 1;
```

Output: DUMP get\_top\_emp;

```

2017-12-10 23:30:01,466 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 23:30:01,466 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(110,Priyanka,2000,5,110,400)
grunt> DESCRIBE get_top_emp;
get_top_emp: {employee_details::emp_id: int,employee_details::emp_name: chararray,employee_details::emp_salary: int,employee_details::emp_rating: int,employee_expenses::emp_id: int,employee_expenses::emp_expense: int}
grunt>

```

acadgild@localhost: ~

Step 7: Get employee id and employee name from above relation.

```
emp_with_max_expense_final = FOREACH get_top_emp GENERATE employee_details::emp_id, employee_details::emp_name;
```

Output: DUMP emp\_with\_max\_expense\_final;

```

2017-12-10 23:30:51,238 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 23:30:51,239 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(110,Priyanka)
grunt> DESCRIBE emp_with_max_expense_final;
emp_with_max_expense_final: {employee_details::emp_id: int,employee_details::emp_name: chararray}
grunt>

```

acadgild@localhost: ~

**Task 4:** List of employees (employee id and employee name) having entries in employee\_expenses file.

**Solution:**

Step 1: Join relations employee\_details and employee\_expenses since we require data from both.

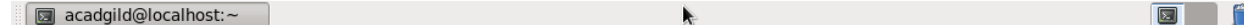
```
joined_emp_data = JOIN employee_details BY emp_id, employee_expenses BY emp_id;Output:
```

Output: DUMP joined\_emp\_data;

```

2017-12-10 23:39:51,334 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 23:39:51,334 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh,20000,1,101,100)
(101,Amitabh,20000,1,101,200)
(102,Shahrukh,10000,2,102,400)
(102,Shahrukh,10000,2,102,100)
(104,Anubhav,5000,4,104,300)
(105,Pawan,2500,5,105,100)
(110,Priyanka,2000,5,110,400)
(114,Madhuri,2000,2,114,200)
grunt> DESCRIBE joined_emp_data;
joined_emp_data: {employee_details::emp_id: int,employee_details::emp_name: chararray,employee_details::emp_salary: int,employee_details::emp_rating: int,employee_expenses::emp_id: int,employee_expenses::emp_expense: int}
grunt>

```



Step 2: Get all the employee names and their id's from joined relation.

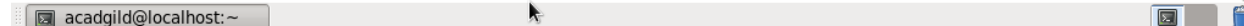
```
emp_info = FOREACH joined_emp_data GENERATE employee_details::emp_id,
employee_details::emp_name;
```

Output: DUMP emp\_info;

```

2017-12-10 23:53:52,127 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 23:53:52,127 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh)
(101,Amitabh)
(102,Shahrukh)
(102,Shahrukh)
(104,Anubhav)
(105,Pawan)
(110,Priyanka)
(114,Madhuri)
grunt> DESCRIBE emp_info;
emp_info: {employee_details::emp_id: int,employee_details::emp_name: chararray}
grunt>

```



**Note:** There is a tuple in employee\_expenses with emp\_id = 119 which is not present in employee\_details. So I have applied normal join (inner join) on source datasets.

Step 3: As we can see in the above output, there are some duplicate tuples. We need to apply DISTINCT operation to remove duplicates.

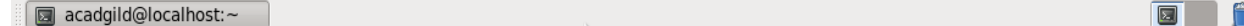
```
emp_info_final = DISTINCT emp_info;
```

Output: DUMP emp\_info\_final;

```

2017-12-10 23:52:27,363 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-10 23:52:27,363 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh)
(102,Shahrukh)
(104,Anubhav)
(105,Pawan)
(110,Priyanka)
(114,Madhuri)
grunt> DESCRIBE emp_info_final;
emp_info_final: {employee_details::emp_id: int,employee_details::emp_name: chararray}
grunt>

```



**Task 5:** List of employees (employee id and employee name) having no entry in employee\_expenses file.

**Solution:**

Step 1: Apply left outer join operation on employee\_details and employee\_expenses since we require all tuples from the former and none from the latter.

```
left_join_emp_details = JOIN employee_details BY emp_id LEFT OUTER, employee_expenses BY emp_id;
```

Output: DUMP left\_join\_emp\_details;

```
2017-12-11 01:04:20,043 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-11 01:04:20,043 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh,20000,1,101,100)
(101,Amitabh,20000,1,101,200)
(102,Shahrukh,10000,2,102,400)
(102,Shahrukh,10000,2,102,100)
(103,Akshay,11000,3,,)
(104,Anubhav,5000,4,104,300)
(105,Pawan,2500,5,105,100)
(106,Aamir,25000,1,,)
(107,Salman,17500,2,,)
(108,Ranbir,14000,3,,)
(109,Katrina,1000,4,,)
(110,Priyanka,2000,5,110,400)
(111,Tushar,500,1,,)
(112,Ajay,5000,2,,)
(113,Jubeen,1000,1,,)
(114,Madhuri,2000,2,114,200)
grunt> DESCRIBE left_join_emp_details;
left_join_emp_details: {employee_details::emp_id: int,employee_details::emp_name: chararray,employee_details::emp_salary: int,employee_details::emp_rating: int,employee_expenses::emp_id: int,employee_expenses::emp_expense: int}
grunt>
```

Step 2: Filter out those records for which emp\_id from employee\_expenses relation is null.

```
joined_emp_details = FILTER left_join_emp_details BY employee_expenses::emp_id IS NULL;
```

Output: DUMP joined\_emp\_details;

```
2017-12-11 01:05:58,346 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2017-12-11 01:05:58,346 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(103,Akshay,11000,3,,)
(106,Aamir,25000,1,,)
(107,Salman,17500,2,,)
(108,Ranbir,14000,3,,)
(109,Katrina,1000,4,,)
(111,Tushar,500,1,,)
(112,Ajay,5000,2,,)
(113,Jubeen,1000,1,,)
grunt> DESCRIBE joined_emp_details;
joined_emp_details: {employee_details::emp_id: int,employee_details::emp_name: chararray,employee_details::emp_salary: int,employee_details::emp_rating: int,employee_expenses::emp_id: int,employee_expenses::emp_expense: int}
grunt>
```



Step 3: Get employee id and employee name from above relation.

```
employees_from_emp_details= FOREACH joined_emp_details GENERATE employee_details::  
emp_id, employee_details::emp_name;
```

Output: DUMP employees\_from\_emp\_details;

```
2017-12-11 01:06:41,353 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1  
2017-12-11 01:06:41,354 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to pro  
cess : 1  
(103,Akshay)  
(106,Aamir)  
(107,Salman)  
(108,Ranbir)  
(109,Katrina)  
(111,Tushar)  
(112,Ajay)  
(113,Jubeen)  
grunt> DESCRIBE employees_from_emp_details;  
employees_from_emp_details: {employee_details::emp_id: int,employee_details::emp_name: chararray}  
grunt>
```

acadmild@localhost:~